Exploratory sensitivity · n=1,048,576, m/n=16

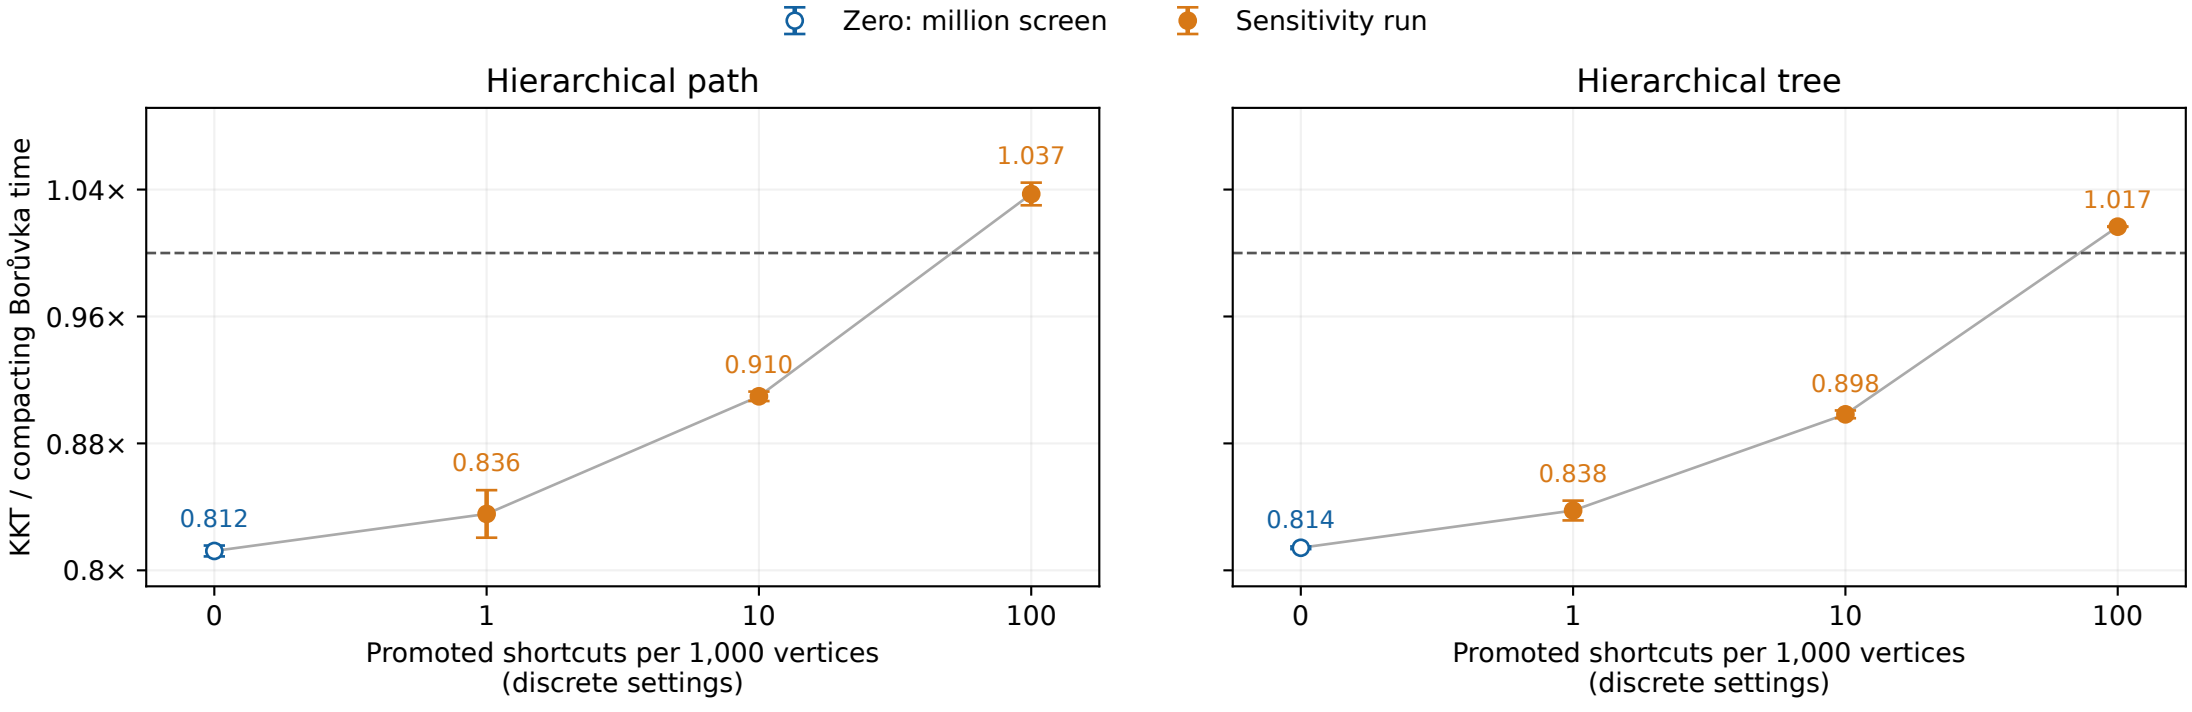

Each point uses matched KKT/compaction calls with graph seeds 17,42, sampling seeds 0,1 and two repetitions; whiskers span two graph medians, not confidence intervals. Zero uses the same seeds and topology from the separate million screen.
That baseline capture includes a logged 1,030 s clamshell sleep, without a long recorded solver call; its suspension phase is unknown.
Promotions equal floor(n × setting / 1000), not a fraction of all edges. The campaign is exploratory; it cannot rescue a failed held-out confirmation. Separate runs may have different desktop conditions.