Current KKT complexity diagnostics — Google Benchmark manual clock

Time per represented input item: T / (n + m) · all 11 sizes · no pooled curve or fitted extrapolation

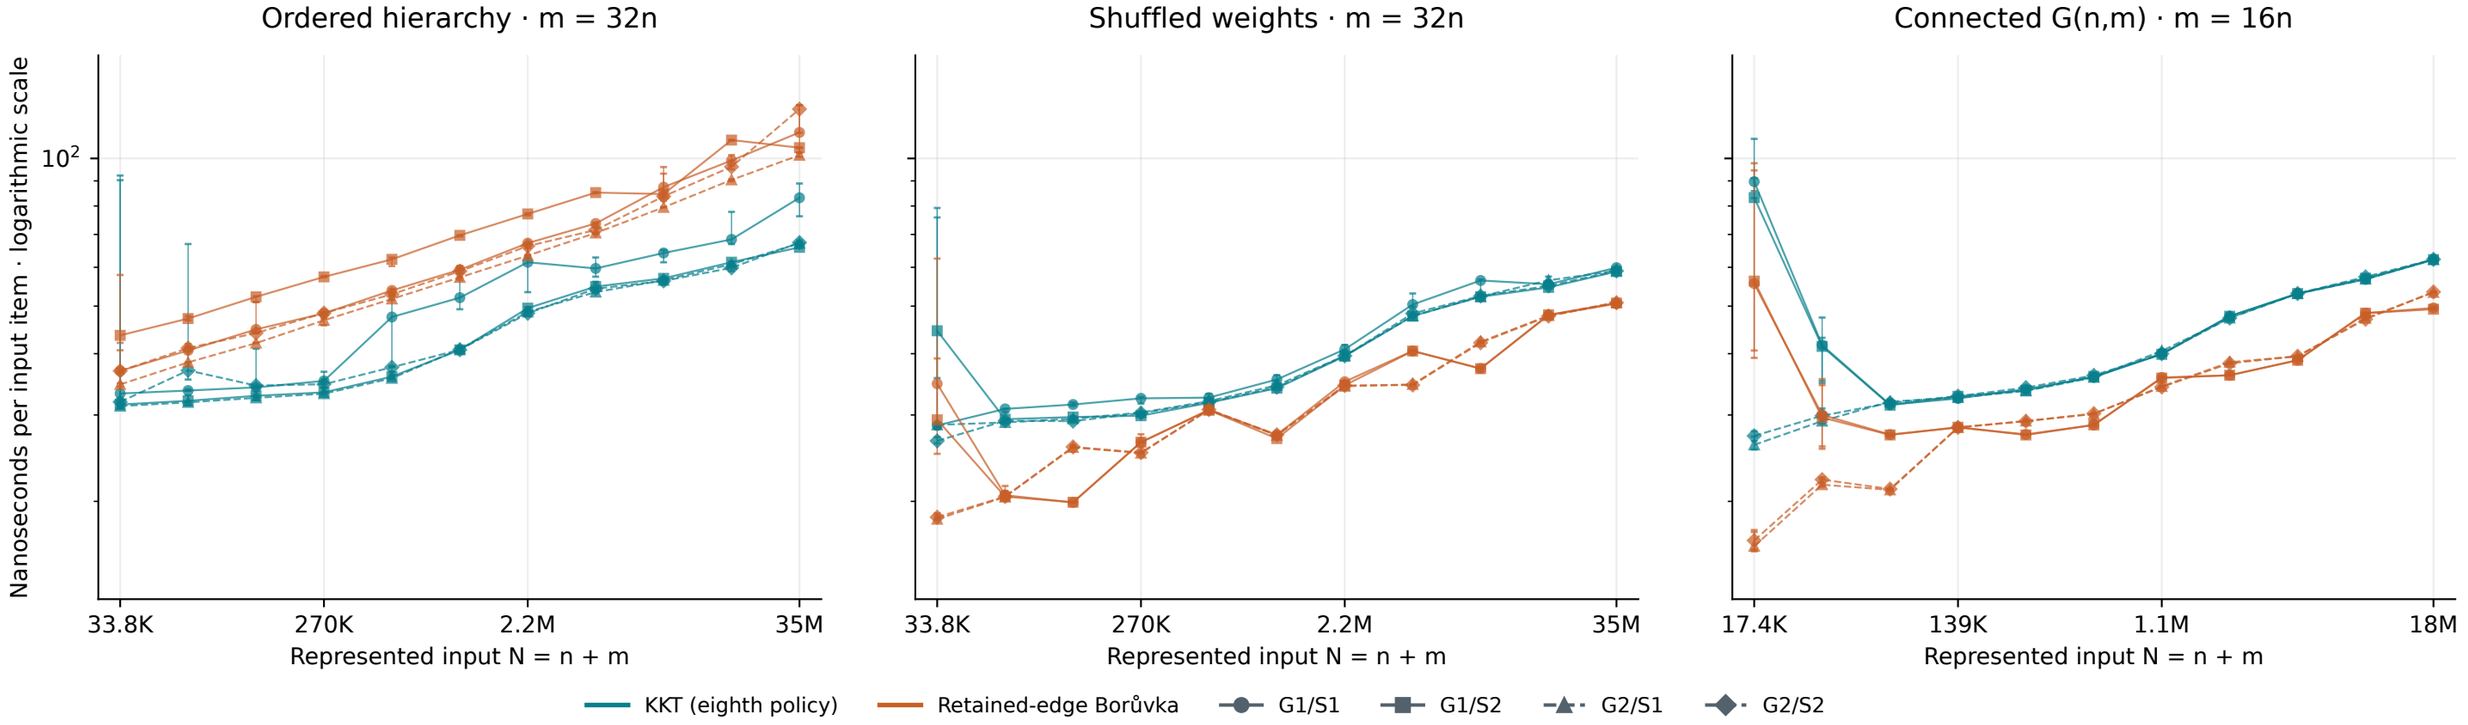

Median and complete min–max of three library two-call means per realization/size. Lines only connect observations; ranges are not confidence intervals.

G1: ascending sizes; G2: descending sizes (graph/order confounded). Two sampler seeds; all flags retained; all four figures share y limits.