Absolute runtime of all six MST implementations
Process CPU time

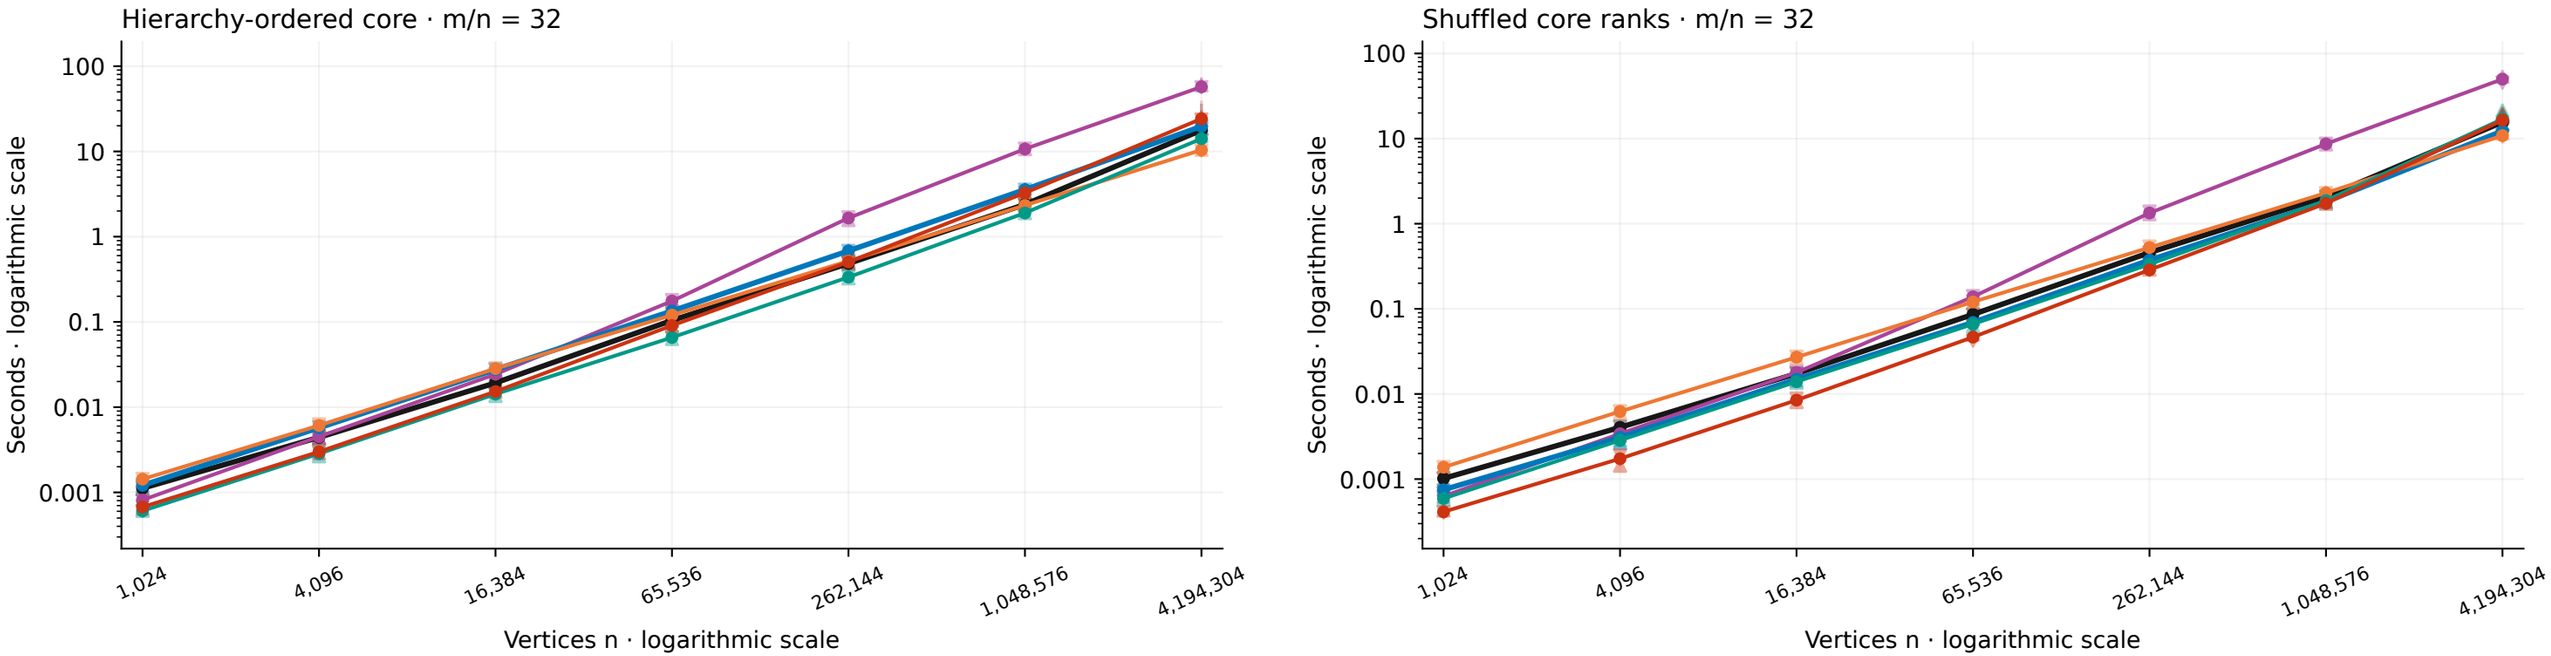

Solid points: medians of twelve durations. Faint triangles: two six-trial graph medians. Vertical spans: every trial-range endpoint.
Two graph seeds × one fixed scalar sampling seed; repetitions reuse each realization. Ranges are descriptive, not confidence intervals.
Lines connect measured sizes, not inferred thresholds. All 0 flagged calls retained. No normalization or historical pooling.