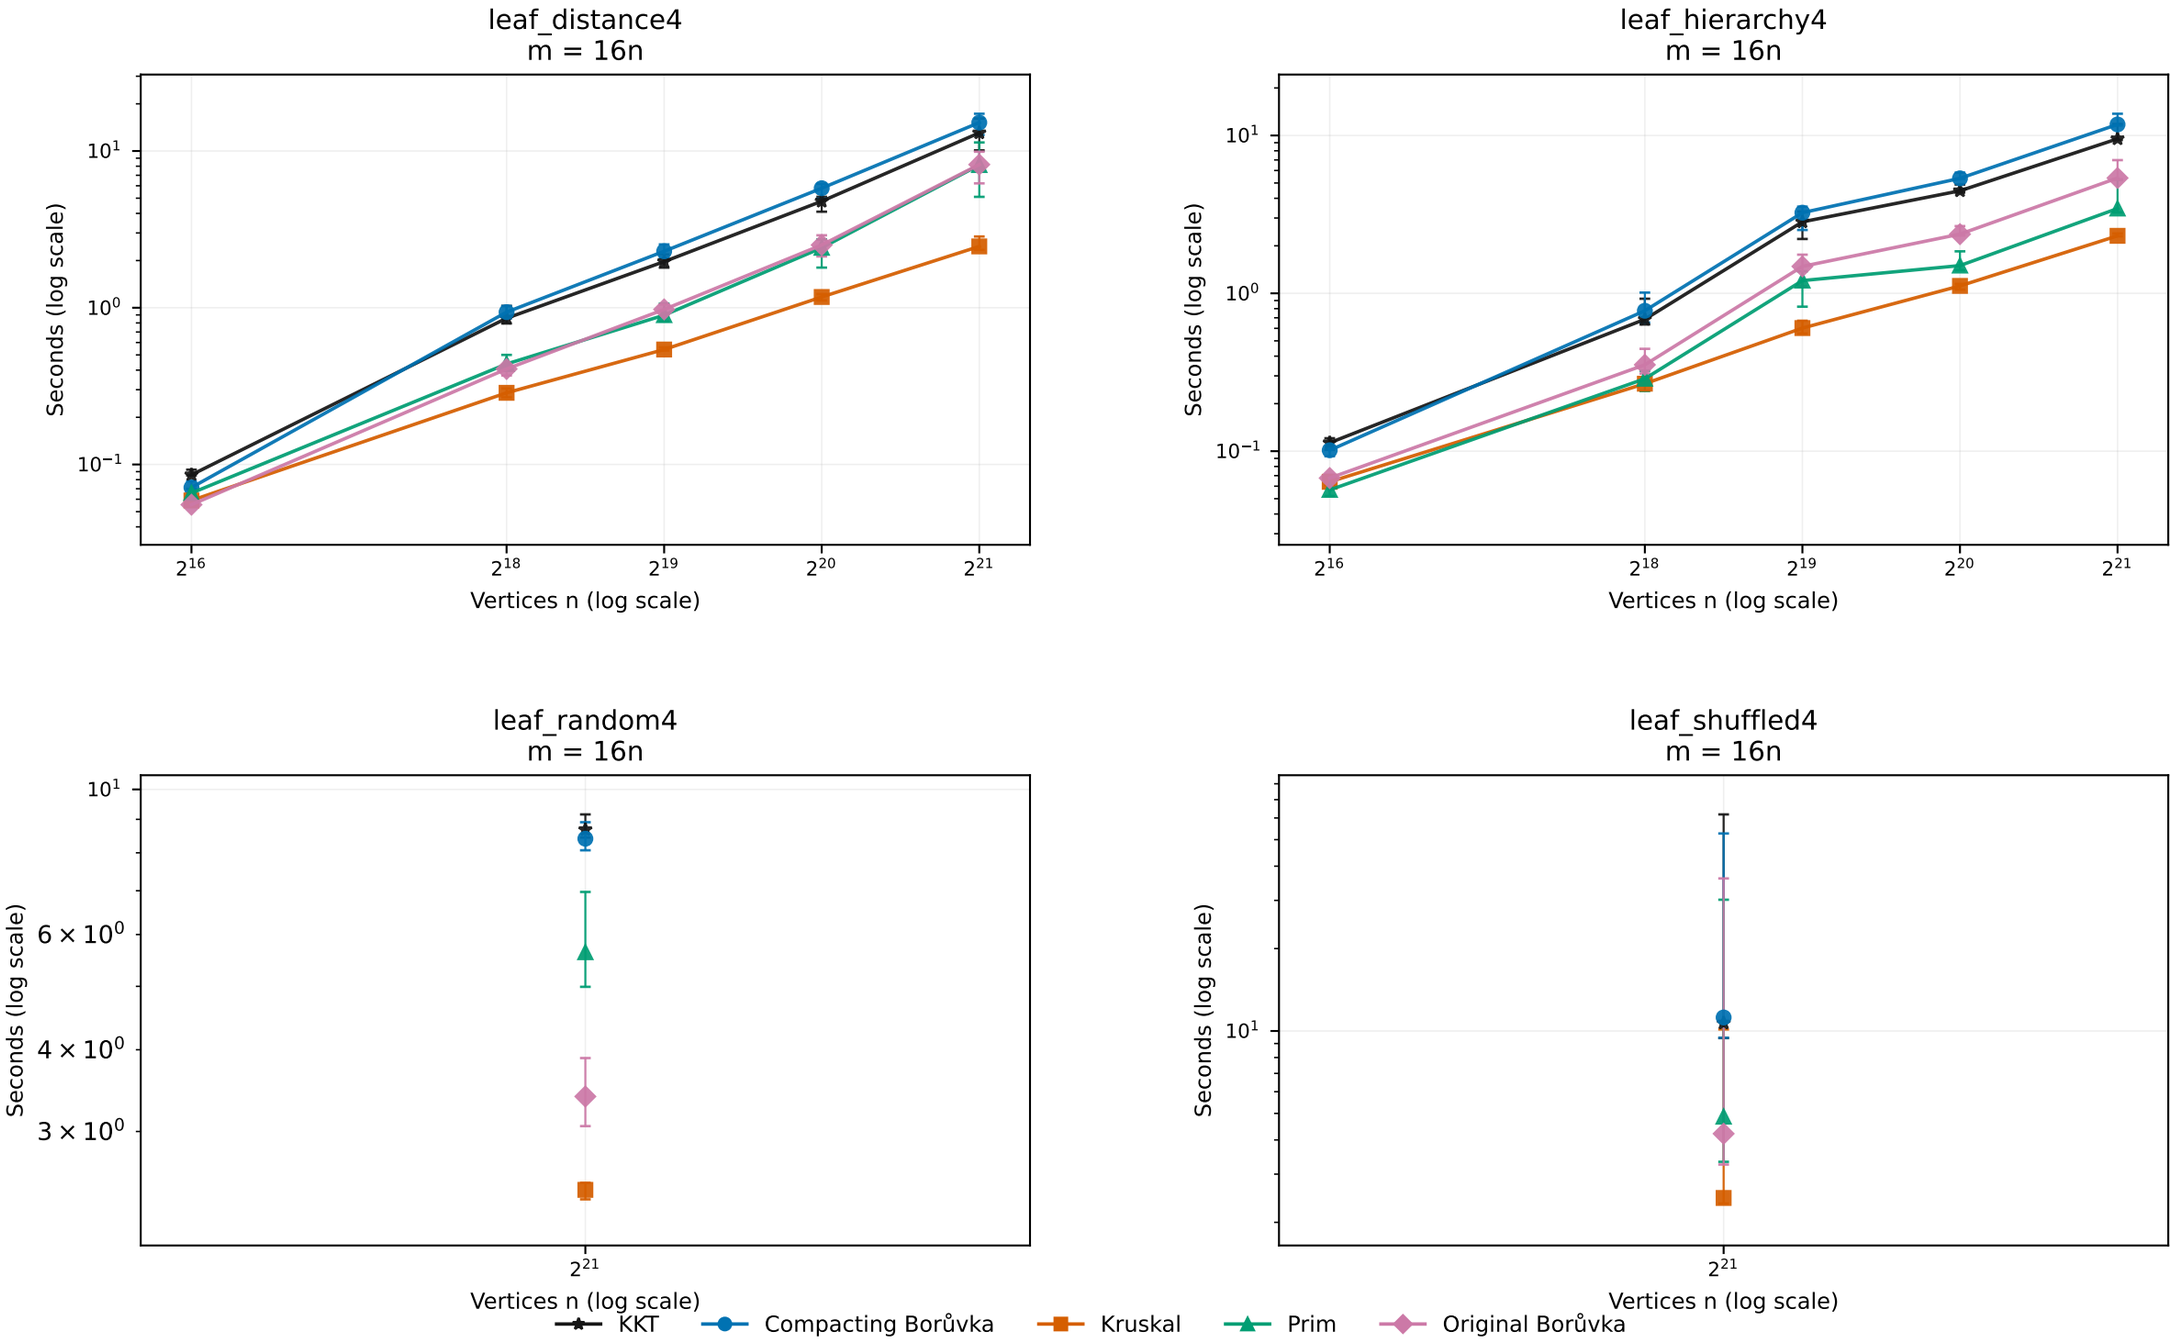
B interrupted + recovery
Process CPU time (secondary) — separate algorithm runtimes
3 graph seed(s) × 2 sampling seed(s) × 5 timed repetitions; 1 checked warmup(s) per algorithm
leaf_distance4
m = 16n
leaf_hierarchy4
m = 16n
leaf_random4
m = 16n
leaf_shuffled4
m = 16n
Seconds (log scale)
Vertices n (log scale)
KKT
Compacting Borůvka
Kruskal
Prim
Original Borůvka
Markers: hierarchical medians. Error bars: observed individual-trial min–max, not confidence intervals.
Lines only connect measured sizes; no smoothing or inferred crossover. Campaigns remain separate.