search-2026-10-13/attachment-large-analysis
Process CPU time (secondary) — paired KKT ratios; below 1 favors KKT
1 graph seed(s) × 1 sampling seed(s) × 5 timed repetitions; 1 checked warmup(s) per algorithm

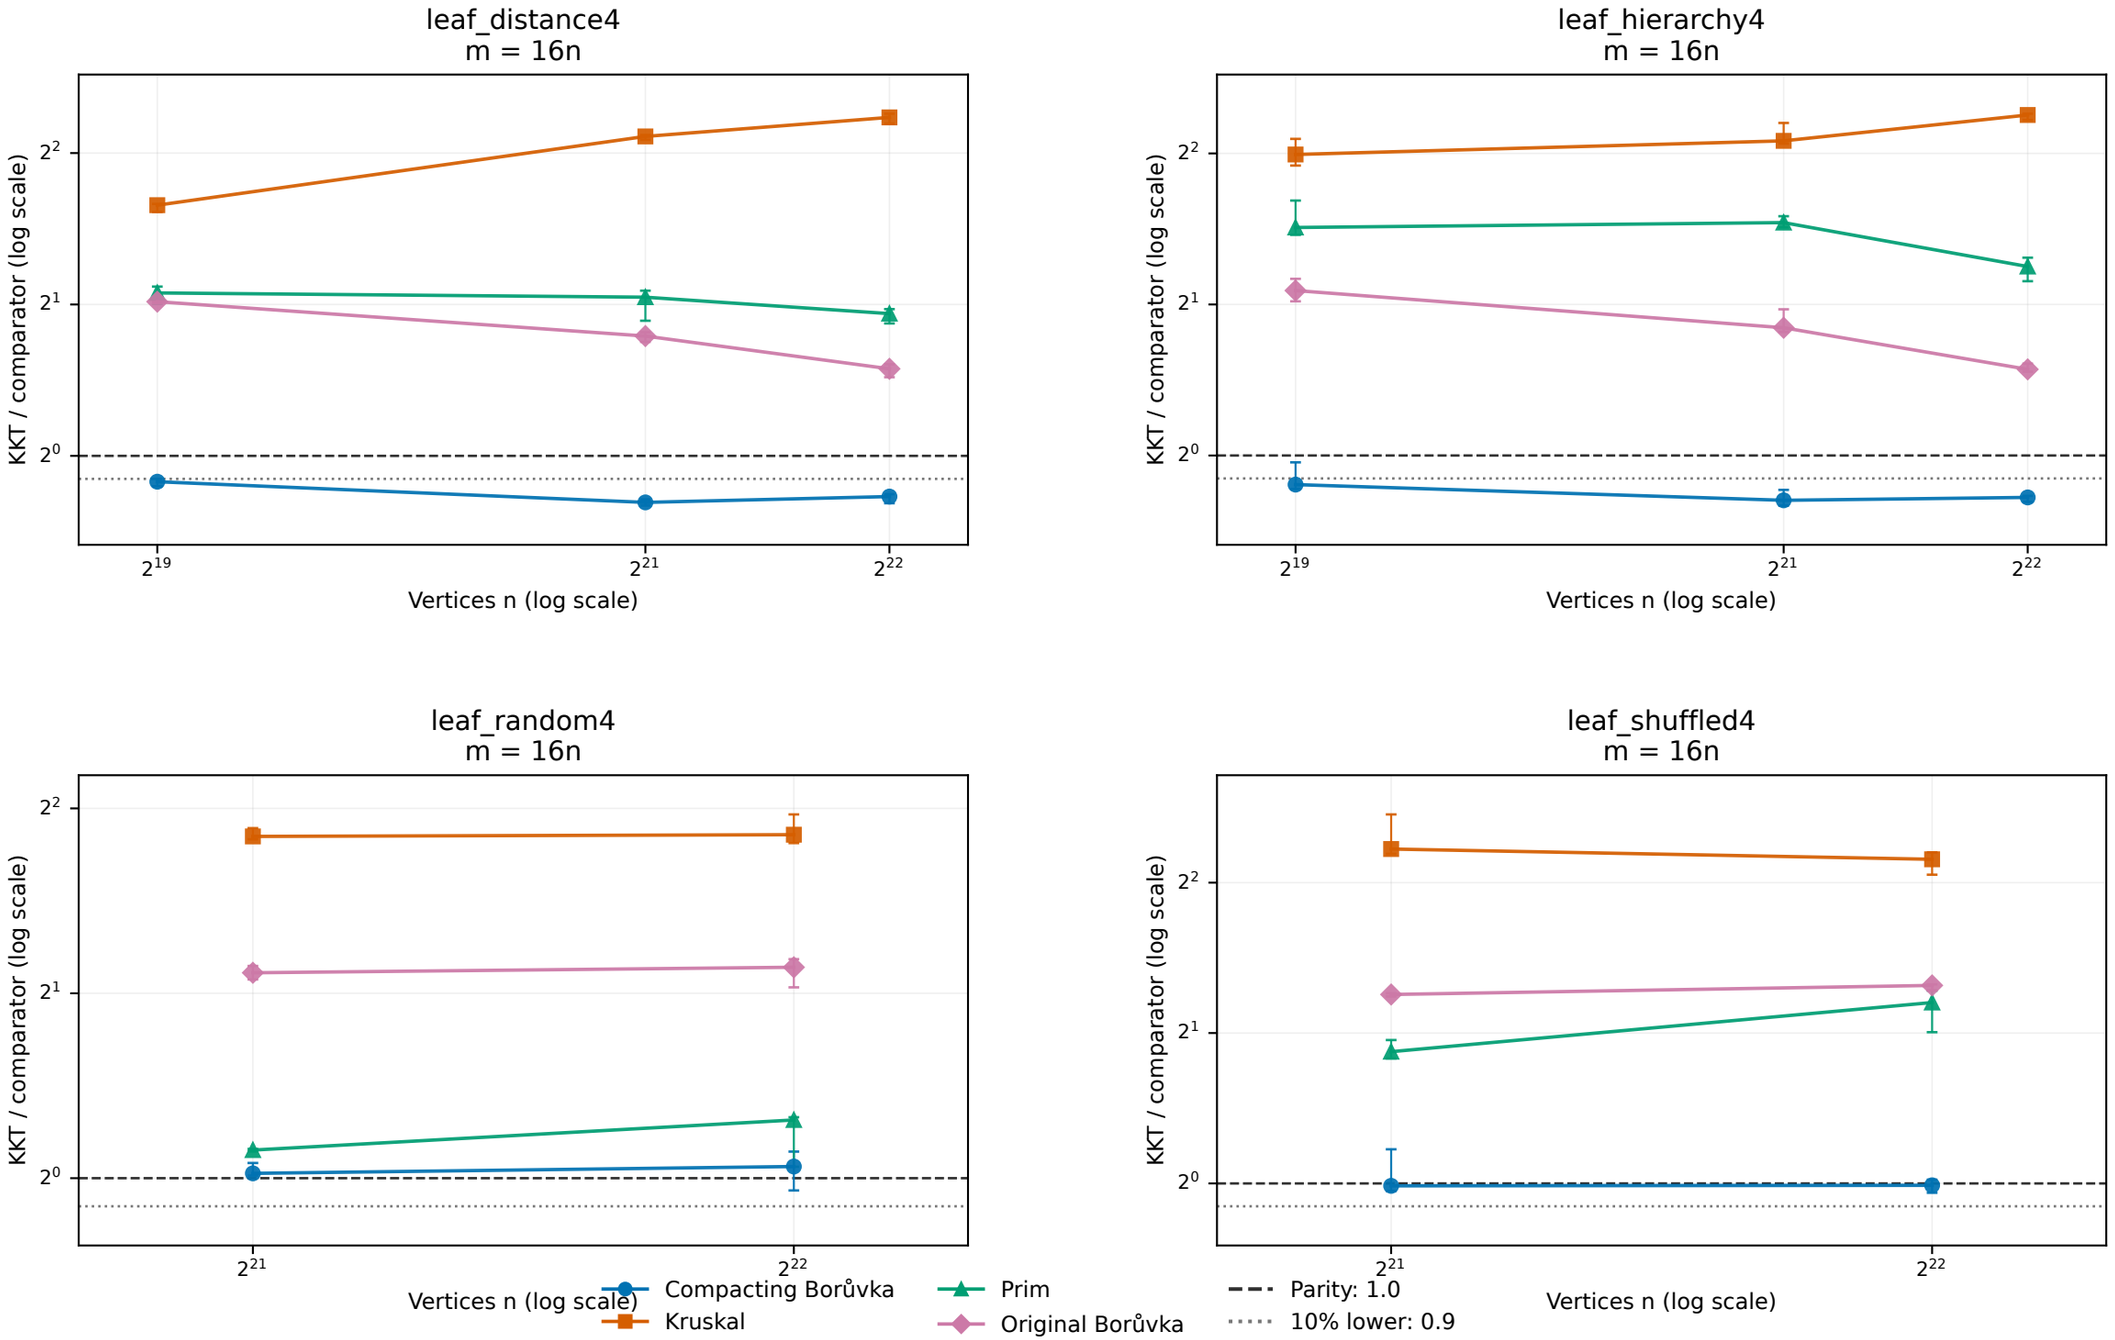

Markers: hierarchical medians. Error bars: observed individual-trial min–max, not confidence intervals.
Lines only connect measured sizes; no smoothing or inferred crossover. Campaigns remain separate.